

```

...
Problem Statement -- Implementing Feedforward neural networks with
Keras and TensorFlow
a. Import the necessary packages
b. Load the training and testing data (MNIST/CIFAR10)
c. Define the network architecture using Keras
d. Train the model using SGD
e. Evaluate the network
f. Plot the training loss and accuracy
...

```

```

# a. importing packages
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import random

# b. LOAD THE TRAINING AND TESTING DATA (MNIST)
mnist = tf.keras.datasets.mnist
(x_train,y_train),(x_test,y_test) = mnist.load_data()

x_train = x_train/255
x_test = x_test/255

Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/mnist.npz
11490434/11490434 ━━━━━━━━━━ 0s 0us/step

# c. DEFINE THE NETWORK ARCHITECTURE USING KERAS ->
model =keras.Sequential([
    keras.layers.Flatten(input_shape=(28,28)),
    keras.layers.Dense(128,activation='relu'),
    keras.layers.Dense(10,activation='softmax')
])

model.summary()

```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/reshaping/  
 flatten.py:37: UserWarning: Do not pass an `input\_shape`/`input\_dim`  
 argument to a layer. When using Sequential models, prefer using an  
 `Input(shape)` object as the first layer in the model instead.  
 super().\_\_init\_\_(\*\*kwargs)

Model: "sequential"

Layer (type) Param #	Output Shape

```

| flatten (Flatten) | (None, 784) |
0 |
|-----|
| dense (Dense) | (None, 128) |
100,480 |
|-----|
| dense_1 (Dense) | (None, 10) |
1,290 |
|-----|
|-----|
```

Total params: 101,770 (397.54 KB)

Trainable params: 101,770 (397.54 KB)

Non-trainable params: 0 (0.00 B)

```
# d. TRAIN THE MODEL USING SGD
model.compile(optimizer='sgd',
  loss='sparse_categorical_crossentropy',
  metrics=['accuracy'])
)

history=model.fit(x_train,y_train,validation_data=(x_test,y_test),epoch
hs=3)

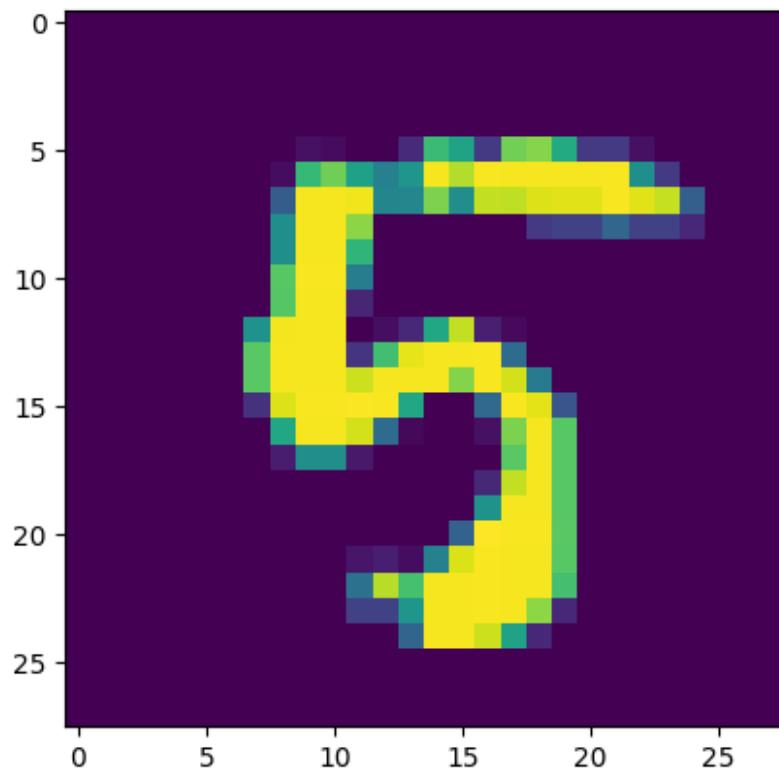
Epoch 1/3
1875/1875 7s 4ms/step - accuracy: 0.7399 - loss:
1.0029 - val_accuracy: 0.9069 - val_loss: 0.3518
Epoch 2/3
1875/1875 5s 3ms/step - accuracy: 0.9053 - loss:
0.3476 - val_accuracy: 0.9197 - val_loss: 0.2899
Epoch 3/3
1875/1875 10s 3ms/step - accuracy: 0.9189 - loss:
0.2963 - val_accuracy: 0.9298 - val_loss: 0.2586

# e. EVALUATE THE NETWORK
test_loss,test_acc = model.evaluate(x_test,y_test)
print("Loss=% .3f" %test_loss)
print("Accuracy=% .3f" %test_acc)

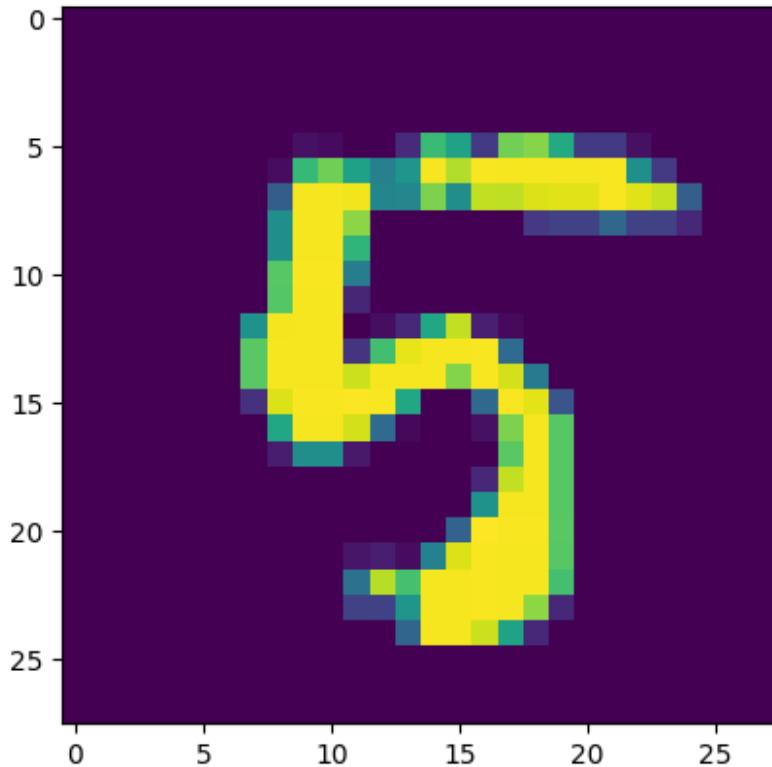
n=random.randint(0,9999)
plt.imshow(x_test[n])
plt.show()
predicted_value = model.predict(x_test)
plt.imshow(x_test[n])
plt.show()

print("Predicted Value:",predicted_value[n])
```

313/313 ━━━━━━━━ 1s 2ms/step - accuracy: 0.9193 - loss:  
0.2955  
Loss=0.259  
Accuracy=0.930

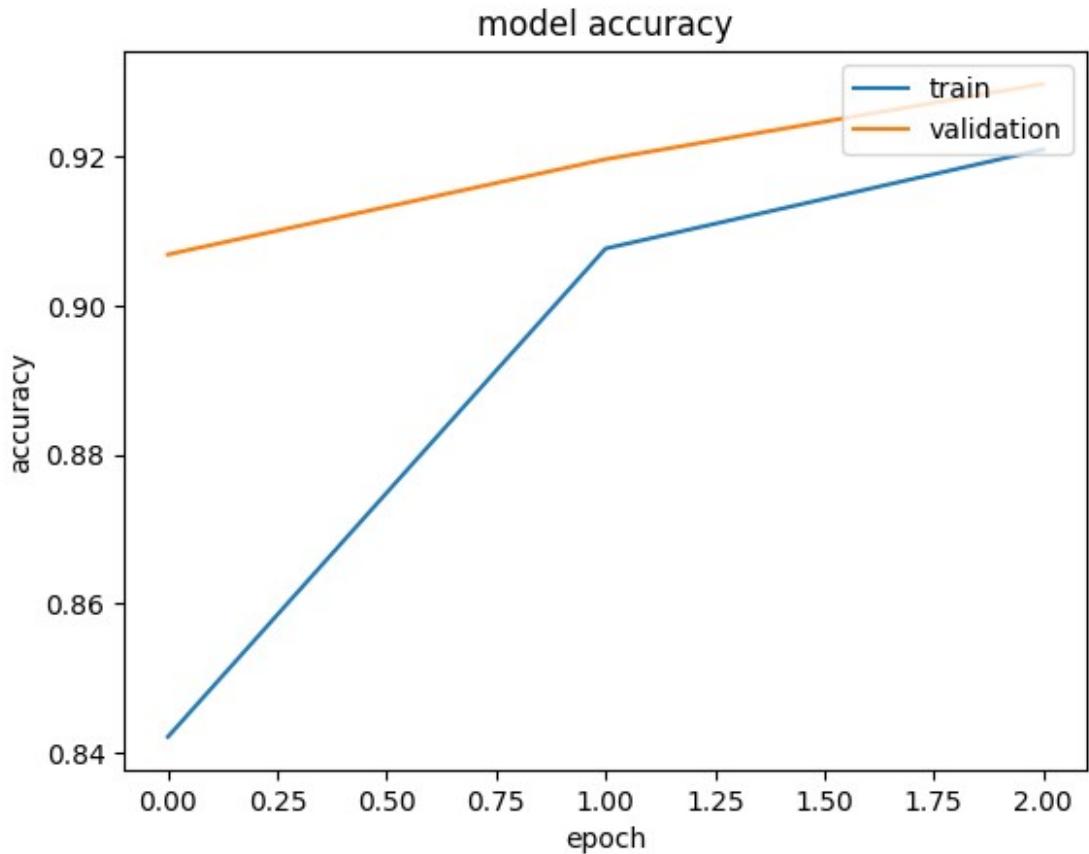


313/313 ━━━━━━━━ 1s 1ms/step



```
Predicted Value: [5.0470990e-04 3.6435162e-07 5.7552425e-05  
2.9780556e-04 1.9890675e-02  
9.0380585e-01 2.3467431e-03 9.1786347e-07 7.1932137e-02 1.1632276e-  
03]
```

```
# f. PLOT THE TRAINING LOSS AND ACCURACY  
    # plotting the training accuracy  
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('model accuracy')  
plt.ylabel('accuracy')  
plt.xlabel('epoch')  
plt.legend(['train', 'validation'], loc='upper right')  
plt.show()
```



```
# plotting the training loss  
  
plt.plot(history.history["loss"])  
plt.plot(history.history["val_loss"])  
plt.title("model loss")  
plt.ylabel("loss")  
plt.xlabel("epoch")  
plt.legend(["train", "validation"], loc="upper left")  
plt.show()
```

