



```
In [1]: """
Problem Statement -- Implementing Feedforward neural networks with Keras and T
a. Import the necessary packages
b. Load the training and testing data (MNIST/CIFAR10)
c. Define the network architecture using Keras
d. Train the model using SGD
e. Evaluate the network
f. Plot the training loss and accuracy
"""
```

```
In [7]: # a. importing packages
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import random
```

```
In [8]: # b. LOAD THE TRAINING AND TESTING DATA (MNIST)
mnist = tf.keras.datasets.mnist
(x_train,y_train),(x_test,y_test) = mnist.load_data()

x_train = x_train/255
x_test = x_test/255
```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 ————— 0s 0us/step

```
In [9]: # c. DEFINE THE NETWORK ARCHITECTURE USING KERAS ->
model =keras.Sequential([
    keras.layers.Flatten(input_shape=(28,28)),
    keras.layers.Dense(128,activation='relu'),
    keras.layers.Dense(10,activation='softmax')
])

model.summary()
```

```
/usr/local/lib/python3.12/dist-packages/keras/src/layers/reshaping/flatten.py:3
7: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. W
hen using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
    super().__init__(**kwargs)
Model: "sequential"
```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100,480
dense_1 (Dense)	(None, 10)	1,290

Total params: 101,770 (397.54 KB)

Trainable params: 101,770 (397.54 KB)

Non-trainable params: 0 (0.00 B)

```
In [10]: # d. TRAIN THE MODEL USING SGD
model.compile(optimizer='sgd',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
)

history=model.fit(x_train,y_train,validation_data=(x_test,y_test),epochs=3)

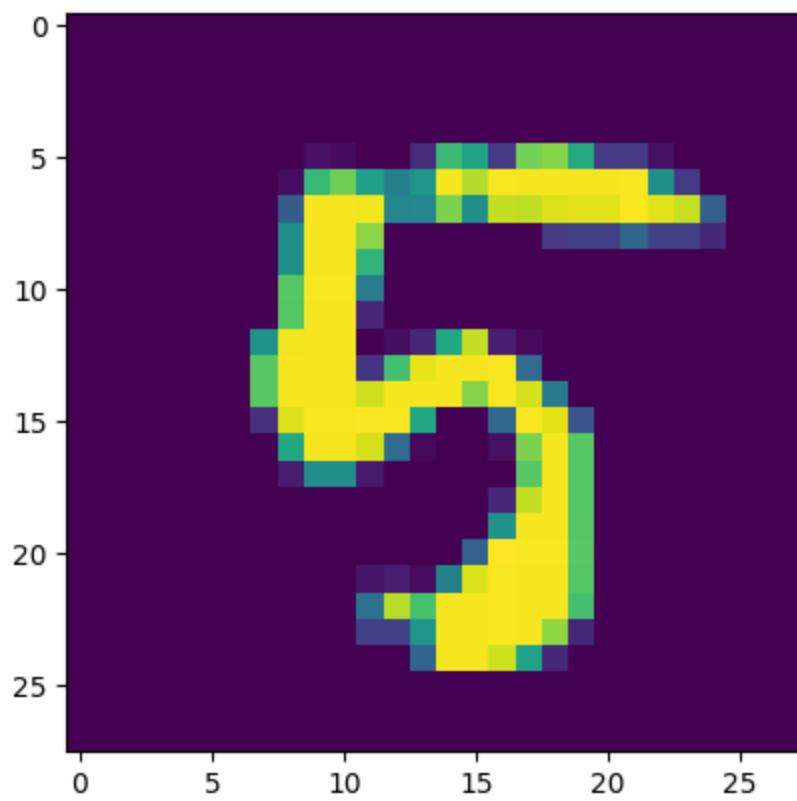
Epoch 1/3
1875/1875 ━━━━━━━━━━ 7s 4ms/step - accuracy: 0.7399 - loss: 1.0029 -
val_accuracy: 0.9069 - val_loss: 0.3518
Epoch 2/3
1875/1875 ━━━━━━━━━━ 5s 3ms/step - accuracy: 0.9053 - loss: 0.3476 -
val_accuracy: 0.9197 - val_loss: 0.2899
Epoch 3/3
1875/1875 ━━━━━━━━━━ 10s 3ms/step - accuracy: 0.9189 - loss: 0.2963 -
val_accuracy: 0.9298 - val_loss: 0.2586
```

```
In [11]: # e. EVALUATE THE NETWORK
test_loss,test_acc = model.evaluate(x_test,y_test)
print("Loss=% .3f" %test_loss)
print("Accuracy=% .3f" %test_acc)

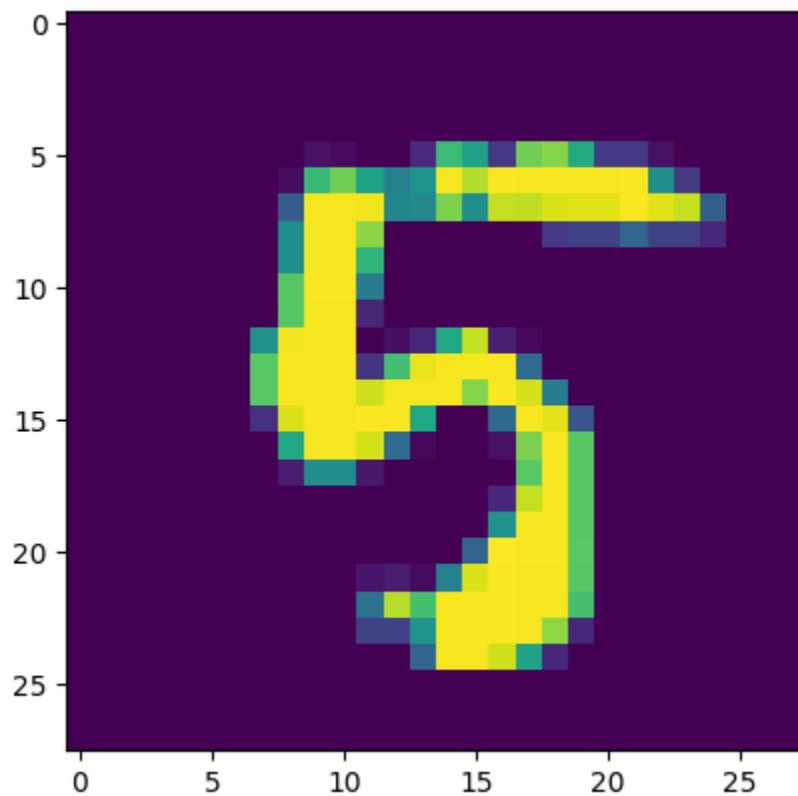
n=random.randint(0,9999)
plt.imshow(x_test[n])
plt.show()
predicted_value = model.predict(x_test)
plt.imshow(x_test[n])
plt.show()

print("Predicted Value:",predicted_value[n])

313/313 ━━━━━━━━━━ 1s 2ms/step - accuracy: 0.9193 - loss: 0.2955
Loss=0.259
Accuracy=0.930
```

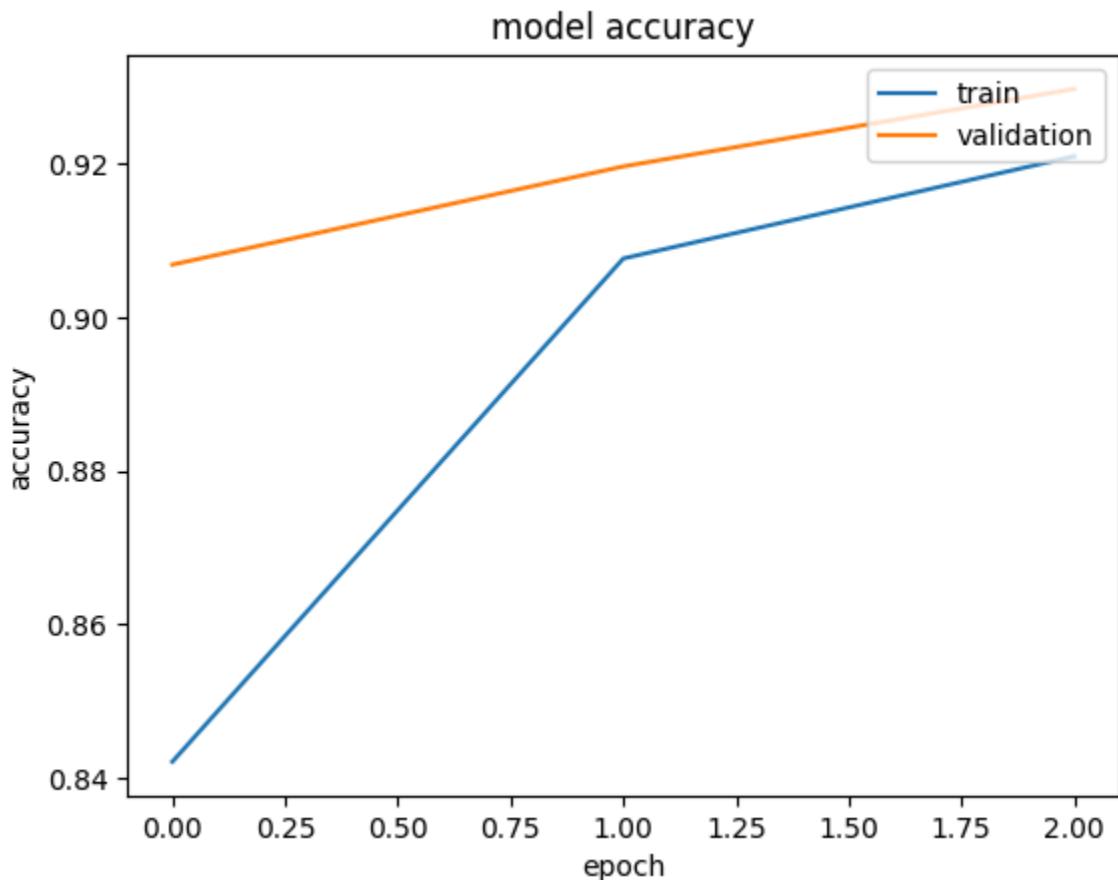


313/313 ━━━━━━ 1s 1ms/step

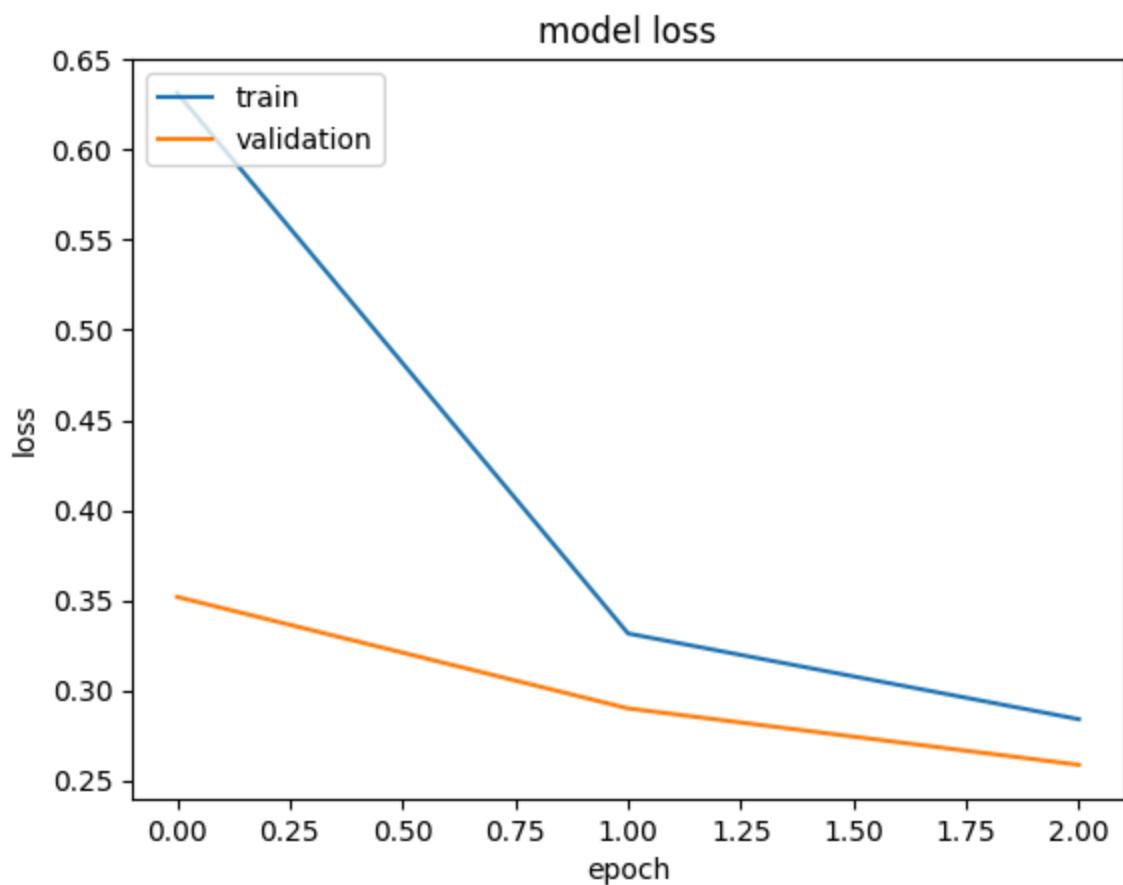


Predicted Value: [5.0470990e-04 3.6435162e-07 5.7552425e-05 2.9780556e-04 1.989
0675e-02
9.0380585e-01 2.3467431e-03 9.1786347e-07 7.1932137e-02 1.1632276e-03]

```
In [12]: # f. PLOT THE TRAINING LOSS AND ACCURACY
    # plotting the training accuracy
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'validation'], loc='upper right')
    plt.show()
```



```
In [14]: # plotting the training loss
plt.plot(history.history["loss"])
plt.plot(history.history["val_loss"])
plt.title("model loss")
plt.ylabel("loss")
plt.xlabel("epoch")
plt.legend(["train", "validation"], loc="upper left")
plt.show()
```



In []:

In []: