

```
...  
Problem Statement - Build the Image classification model by dividing  
the model into following 4 stages:
```

- a. Loading and preprocessing the image data
- b. Defining the model's architecture
- c. Training the model
- d. Estimating the model's performance

```
# IMPORT THE REQUIRED PACKAGES  
import tensorflow as tf  
from keras.models import Sequential  
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D  
import matplotlib.pyplot as plt  
import numpy as np  
  
# a. LOADING AND PRE-PROCESSING THE IMAGE DATA  
mnist = tf.keras.datasets.mnist  
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
input_shape = (28, 28, 1)  
  
Downloading data from https://storage.googleapis.com/tensorflow/tf-  
keras-datasets/mnist.npz  
11490434/11490434 ━━━━━━━━━━ 0s 0us/step  
  
# Making sure that the values are float so that we can get decimal  
points after division  
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)  
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)  
  
print("Data Type of x_train: ", x_train.dtype)  
x_train = x_train.astype('float32')  
x_test = x_test.astype('float32')  
  
Data Type of x_train: float32  
  
# Normalizing the RGB codes by dividing it to the max RGB value.  
x_train /= 255  
x_test /= 255  
print("Shape of Training: ", x_train.shape)  
print("Shape of Testing", x_test.shape)  
  
Shape of Training: (60000, 28, 28, 1)  
Shape of Testing (10000, 28, 28, 1)  
  
# b. DEFINING THE MODEL'S ARCHITECTURE ->  
model = Sequential()  
model.add(Conv2D(28, kernel_size=(3,3), input_shape = input_shape))  
model.add(MaxPooling2D(pool_size=(2, 2)))  
model.add(Flatten())  
model.add(Dense(200, activation = "relu"))
```

```
model.add(Dropout(0.3))
model.add(Dense(10, activation = "softmax"))
model.summary()

/usr/local/lib/python3.12/dist-packages/keras/src/layers/
convolutional/base_conv.py:113: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in
the model instead.
    super().__init__(activity_regularizer=activity_regularizer,
**kwargs)
```

Model: "sequential"

Layer (type)	Output Shape
Param #	
conv2d (Conv2D) 280	(None, 26, 26, 28)
max_pooling2d (MaxPooling2D) 0	(None, 13, 13, 28)
flatten (Flatten) 0	(None, 4732)
dense (Dense) 946,600	(None, 200)
dropout (Dropout) 0	(None, 200)
dense_1 (Dense) 2,010	(None, 10)

Total params: 948,890 (3.62 MB)

Trainable params: 948,890 (3.62 MB)

Non-trainable params: 0 (0.00 B)

```

# c. TRAINING THE MODEL ->
model.compile(optimizer='adam',
loss='sparse_categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train,y_train, epochs=2)

Epoch 1/2
1875/1875 ━━━━━━━━━━ 44s 23ms/step - accuracy: 0.1065 - 
loss: 2.3035
Epoch 2/2
1875/1875 ━━━━━━━━━━ 82s 23ms/step - accuracy: 0.1115 - 
loss: 2.3016

<keras.src.callbacks.history.History at 0x78a6c8f684d0>


# d. ESTIMATING THE MODEL'S PERFORMANCE ->
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Loss=% .3f" %test_loss)
print("Accuracy=% .3f" %test_acc)

# Showing image at position[] from dataset:
image = x_train[0]
plt.imshow(np.squeeze(image), cmap='gray')
plt.show()

# Predicting the class of image:
image = image.reshape(1, image.shape[0], image.shape[1],
image.shape[2])
predict_model = model.predict([image])
print("Predicted class: {}".format(np.argmax(predict_model)))

313/313 ━━━━━━━━━━ 2s 7ms/step - accuracy: 0.1160 - loss:
2.3011
Loss=2.301
Accuracy=0.113

```



1/1 ━━━━━━━━ 0s 81ms/step

Predicted class: 1

```
/usr/local/lib/python3.12/dist-packages/keras/src/models/
functional.py:241: UserWarning: The structure of `inputs` doesn't
match the expected structure.
Expected: keras_tensor
Received: inputs='Tensor(shape=(1, 28, 28, 1))',)
warnings.warn(msg)
```